


Imlab: Dremel

Andreas Zimmerer

Technische Universität München

Garching, 27. Januar 2020



Overview

- Why Dremel?
- How does Dremel Work?
- Example: Record Assembly
- Challenges
- Benchmarks
- Outlook

Why Dremel?

- It's a good idea to support non-relational data:
Alternative would be normalizing data and joining the parts
-> might be very costly
- Dremel utilizes the benefits of column stores for nested data
- Dremel is highly parallelizable

How does Dremel Work?

Shredding documents into columns:

```

DocId: 10      r1
Links
  Forward: 20
  Forward: 40
  Forward: 60
Name
  Language
    Code: 'en-us'
    Country: 'us'
  Language
    Code: 'en'
  Url: 'http://A'
Name
  Url: 'http://B'
Name
  Language
    Code: 'en-gb'
    Country: 'gb'
  
```



DocId0	r	d	L.Backward	r	d	L.Forward	r	d
10	0	0	NULL	0	1	20	0	2
						40	1	2
						60	1	2
N.L.Code	r	d	N.L.Country	r	d	N.Url	r	d
en-us	0	2	us	0	3	http://A	0	2
en	2	2	NULL	2	2	http://B	1	2
NULL	1	1	NULL	1	1	NULL	1	1
en-gb	1	2	gb	1	3			

How do we compute repetition and definition levels?

How does Dremel Work?

Computing definition and repetition levels:

```
message Document {  
  required int64 DocId;  
  optional group Links {  
    repeated int64 Backward;  
    repeated int64 Forward; }  
  repeated group Name {  
    repeated group Language {  
      required string Code;  
      optional string Country; }  
    optional string Url; }}
```

The *definition level* is always fixed for a field, and is used for specifying which field in the message was NULL.

The definition level of a field is the number of potentially undefined (optional and repeated) fields in the path.

The *repetition level* tells us at which level a field is repeated.

The maximum repetition level of a field is the number of repeated fields in the path.

=> with these two values we can losslessly reconstruct records

How does Dremel Work?

Retrieving a record: (1) Create a FSM with the fields

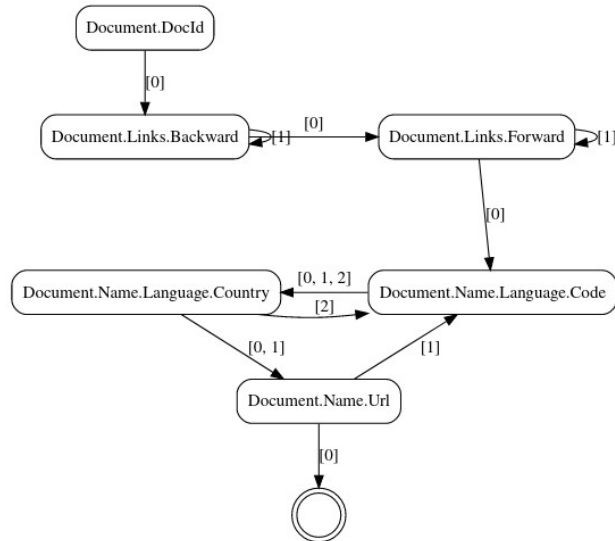
Question: How can we reconstruct a (partial) record?

=> We construct a FSM that defines how we jump between fields.
e.g. given a previous field and a repetition level, what's the next field?

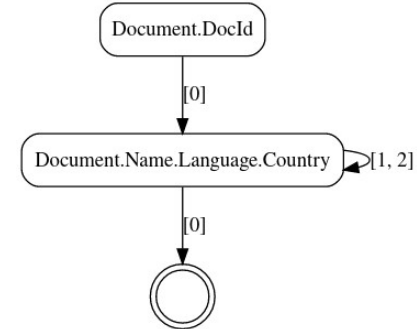
We can also construct a FSM that constructs only a partial record, but still preserves the record structure.

How does Dremel Work?

Full FSM with all fields



Partial FSM for DocId and Country



How does Dremel Work?

Retrieving a record: (2) Assemble record by jumping from field to field with FSM

Example: Record Assembly

Example: Record Assembly

Three Ingredients:

(1) Document Structure

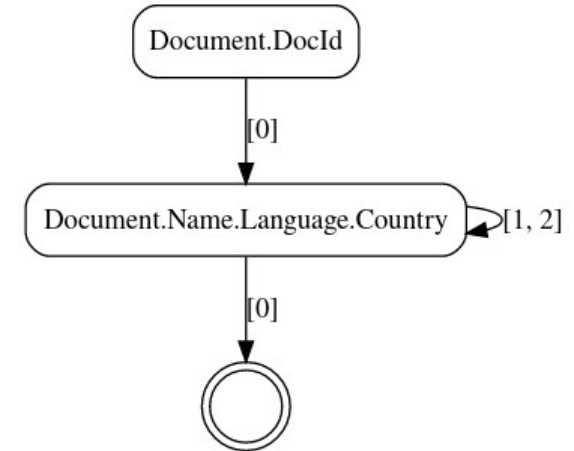
```
message Document {
  required int64 DocId;
  optional group Links {
    repeated int64 Backward;
    repeated int64 Forward; }
  repeated group Name {
    repeated group Language {
      required string Code;
      optional string Country; }
    optional string Url; }}
```

(2) Record Data

DocId	r	d
10	0	0
...	0	0


N.L.Country	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
...	0	3

(3) FSM



Example: Record Assembly

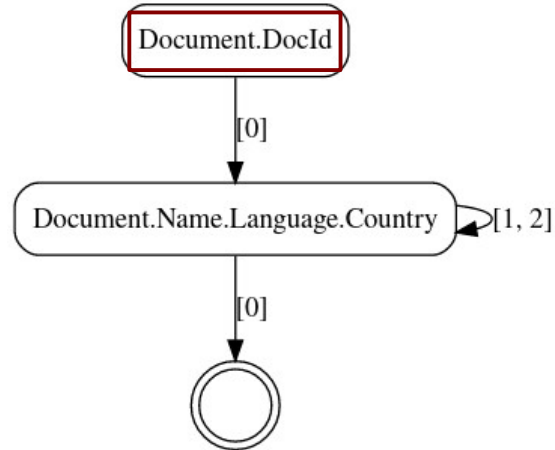
Assembled Record:



DocId	r	d
10	0	0
...	0	0



N.L.Country	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
...	0	3



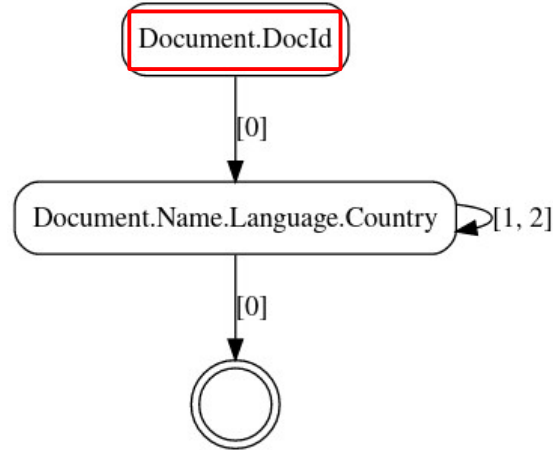
Example: Record Assembly

Assembled Record:

DocId: 10

DocId	r	d
10	0	0
...	0	0

N.L.Country	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
...	0	3



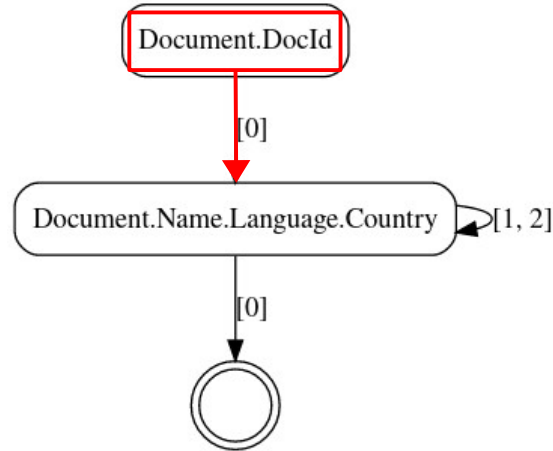
Example: Record Assembly

Assembled Record:

DocId: 10

DocId	r	d
10	0	0
...	0	0

N.L.Country	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
...	0	3



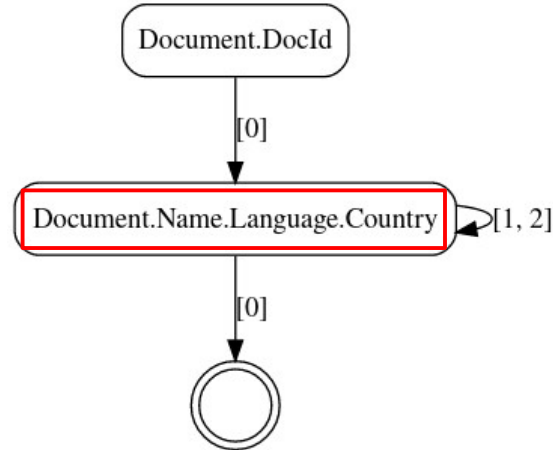
Example: Record Assembly

Assembled Record:

DocId: 10

DocId	r	d
10	0	0
...	0	0

N.L.Country	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
...	0	3



Example: Record Assembly


Assembled Record:

```
DocId: 10
Name: {
  Lang: {
    Country: 'us'
```

DocId	r	d
10	0	0
...	0	0



N.L.Country	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
...	0	3



Example: Record Assembly

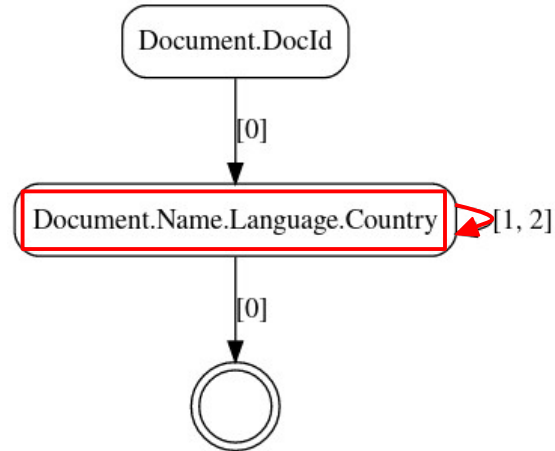
Assembled Record:

```
DocId: 10
Name: {
  Lang: {
    Country: 'us'
  }
}
```

DocId	r	d
10	0	0
...	0	0



N.L.Country	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
...	0	3



Example: Record Assembly


Assembled Record:

```
DocId: 10
Name: {
  Lang: {
    Country: 'us'
  }
}
```

DocId	r	d
10	0	0
...	0	0



N.L.Country	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
...	0	3




Example: Record Assembly

Assembled Record:

```
DocId: 10
Name: {
  Lang: {
    Country: 'us'
  }
  Lang: {
```

DocId	r	d
10	0	0
...	0	0

N.L.Country	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
...	0	3



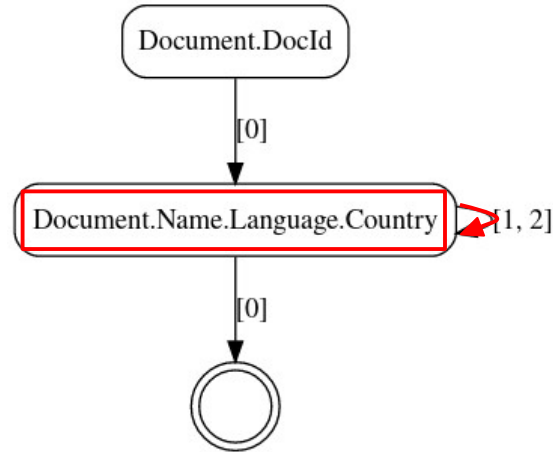
Example: Record Assembly

Assembled Record:

```
DocId: 10
Name: {
  Lang: {
    Country: 'us'
  }
  Lang: { }
```

DocId	r	d
10	0	0
...	0	0

N.L.Country	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
...	0	3



Example: Record Assembly


Assembled Record:

```
DocId: 10
Name: {
  Lang: {
    Country: 'us'
  }
  Lang: { }
```

DocId	r	d
10	0	0
...	0	0



N.L.Country	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
...	0	3




Example: Record Assembly

Assembled Record:

```
DocId: 10
Name: {
  Lang: {
    Country: 'us'
  }
  Lang: { }
```

DocId	r	d
10	0	0
...	0	0

N.L.Country	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
...	0	3



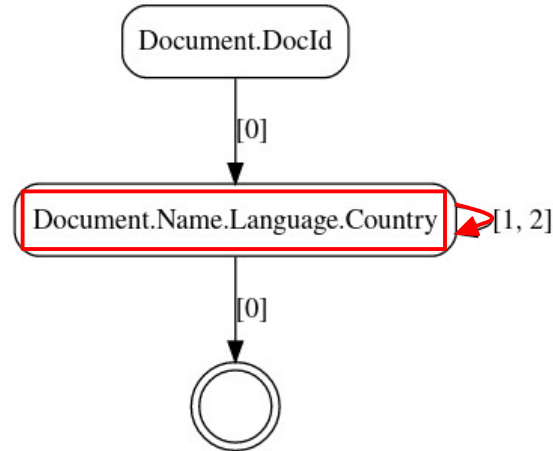
Example: Record Assembly

Assembled Record:

```
DocId: 10
Name: {
  Lang: {
    Country: 'us'
  }
  Lang: {}
}
```

DocId	r	d
10	0	0
...	0	0

N.L.Country	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
...	0	3




Example: Record Assembly

Assembled Record:

```
DocId: 10
Name: {
  Lang: {
    Country: 'us'
  }
  Lang: {}
}
```

DocId	r	d
10	0	0
...	0	0


N.L.Country	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
...	0	3



Example: Record Assembly

Assembled Record:

```
DocId: 10
Name: {
  Lang: {
    Country: 'us'
  }
  Lang: {}
}
Name: {
  Lang: {
    Country: 'gb'
  }
}
```



DocId	r	d
10	0	0
...	0	0



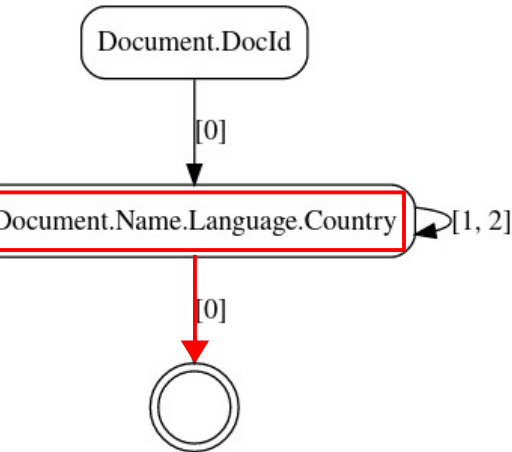
N.L.Country	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
...	0	3



Example: Record Assembly

Assembled Record:

```
DocId: 10
Name: {
  Lang: {
    Country: 'us'
  }
  Lang: {}
}
Name: {
  Lang: {
    Country: 'gb'
  }
}
```



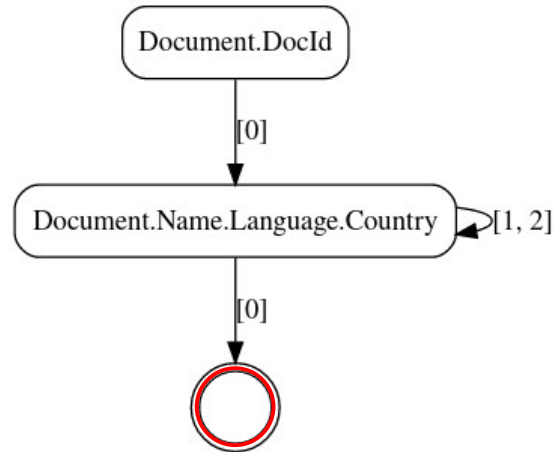
DocId	r	d
10	0	0
...	0	0

N.L.Country	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
...	0	3

Example: Record Assembly

Assembled Record:

```
DocId: 10
Name: {
  Lang: {
    Country: 'us'
  }
  Lang: {}
}
Name: {
  Lang: {
    Country: 'gb'
  }
}
```



DocId	r	d
10	0	0
...	0	0

N.L.Country	r	d
us	0	3
NULL	2	2
NULL	1	1
gb	1	3
...	0	3

Challenges

FSM Construction

```
1 procedure ConstructFSM(Field[] fields):
2   for each field in fields:
3     maxLevel = maximal repetition level of field
4     barrier = next field after field or final FSM state otherwise
5     barrierLevel = common repetition level of field and barrier
6     for each preField before field whose
7       repetition level is larger than barrierLevel:
8       backLevel = common repetition level of preField and field
9       Set transition (field, backLevel) -> preField
10    end for
11    for each level in [barrierLevel+1..maxLevel]
12      that lacks transition from field:
13        Copy transition's destination from that of level-1
14    end for
15    for each level in [0..barrierLevel]:
16      Set transition (field, level) -> barrier
17    end for
18  end for
19 end procedure
```

FSM Construction

```

1 procedure ConstructFSM(Field[] fields):
2   for each field in fields:
3     maxLevel = maximal repetition level of field
4     barrier = next field after field or final FSM state otherwise
5     barrierLevel = common repetition level of field and barrier
6     for each preField before field whose
7       repetition level is larger than barrierLevel:
8       backLevel = common repetition level of preField and field
9       Set transition (field, backLevel) -> preField
10    end for
11    for each level in [barrierLevel+1..maxLevel]
12      that lacks transition from field:
13      Copy transition's destination from that of level-1
14    end for
15    for each level in [0..barrierLevel]:
16      Set transition (field, level) -> barrier
17    end for
18  end for
19 end procedure

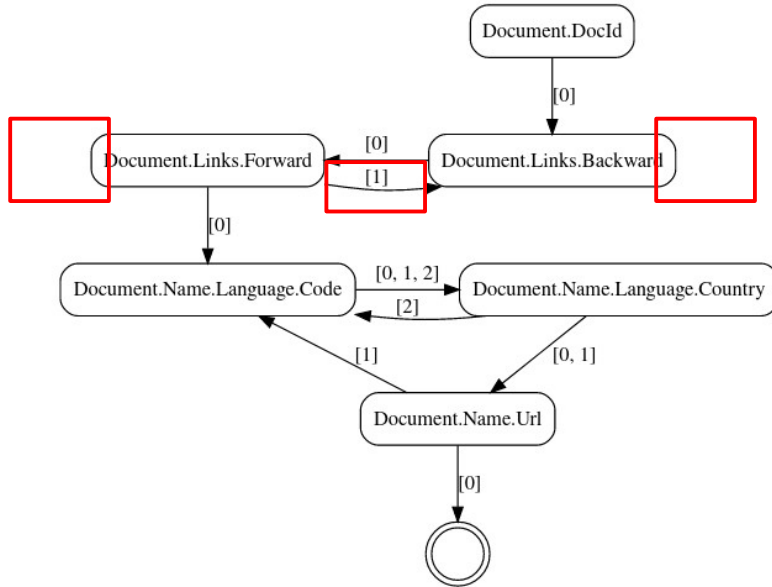
```

No

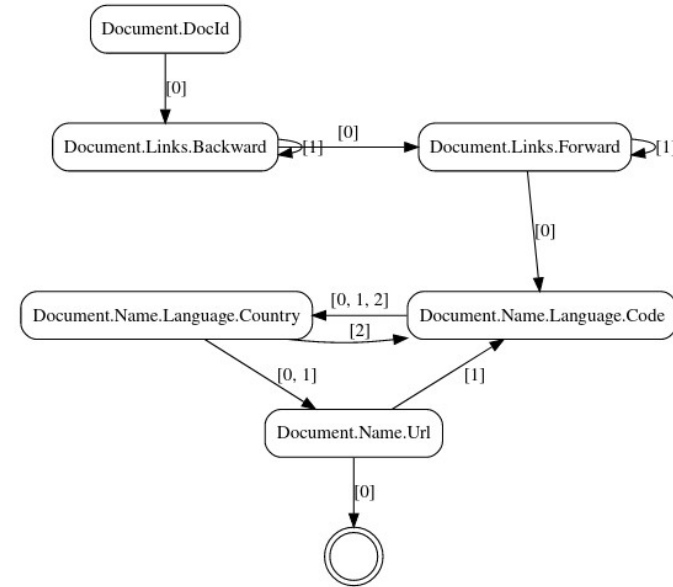
Instead: Insert reflexive edge

FSM Construction: Algorithms Compared

Version from Paper:



Modified Version:



Common Ancestor

```
message Document {  
  required int64 DocId;  
  optional group Links {  
    repeated int64 Backward;  
    repeated int64 Forward; }  
  repeated group Name {  
    repeated group Language {  
      required string Code;  
      optional string Country; }  
    optional string Url; }}  
}
```

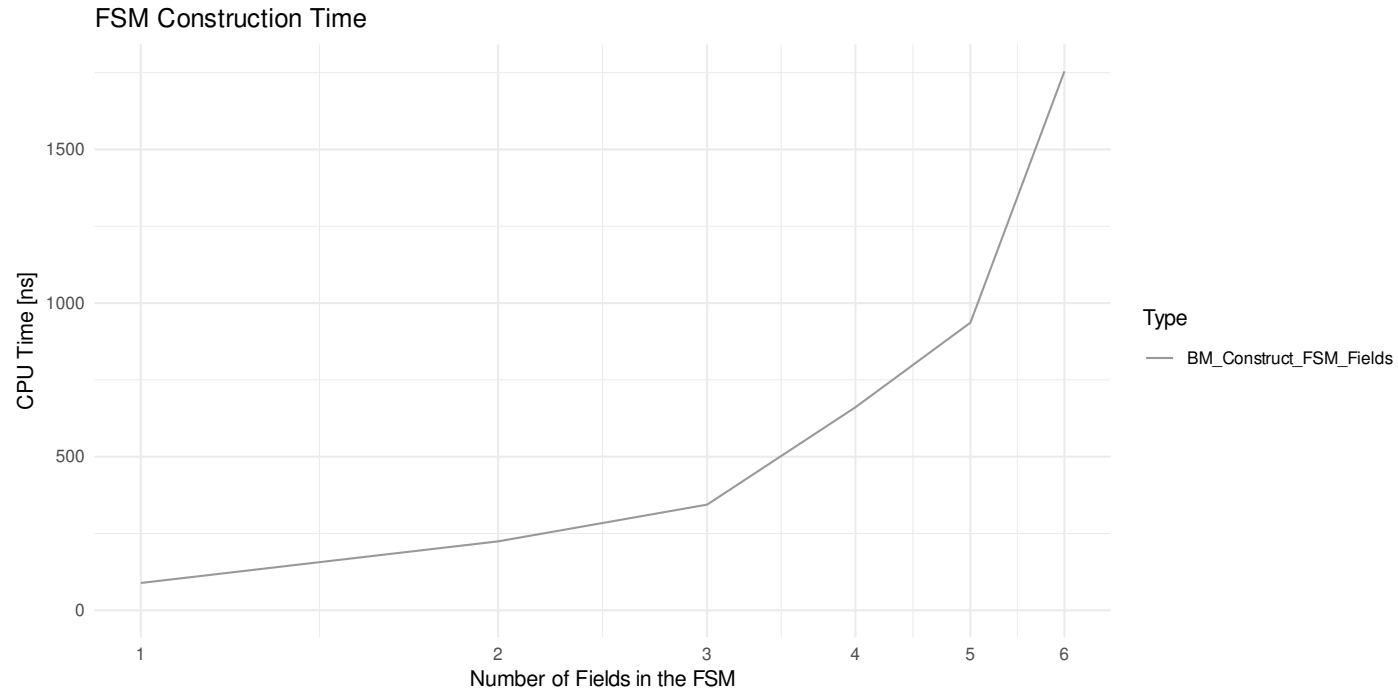
Order is important!

CoAn of Code and Country -> Language

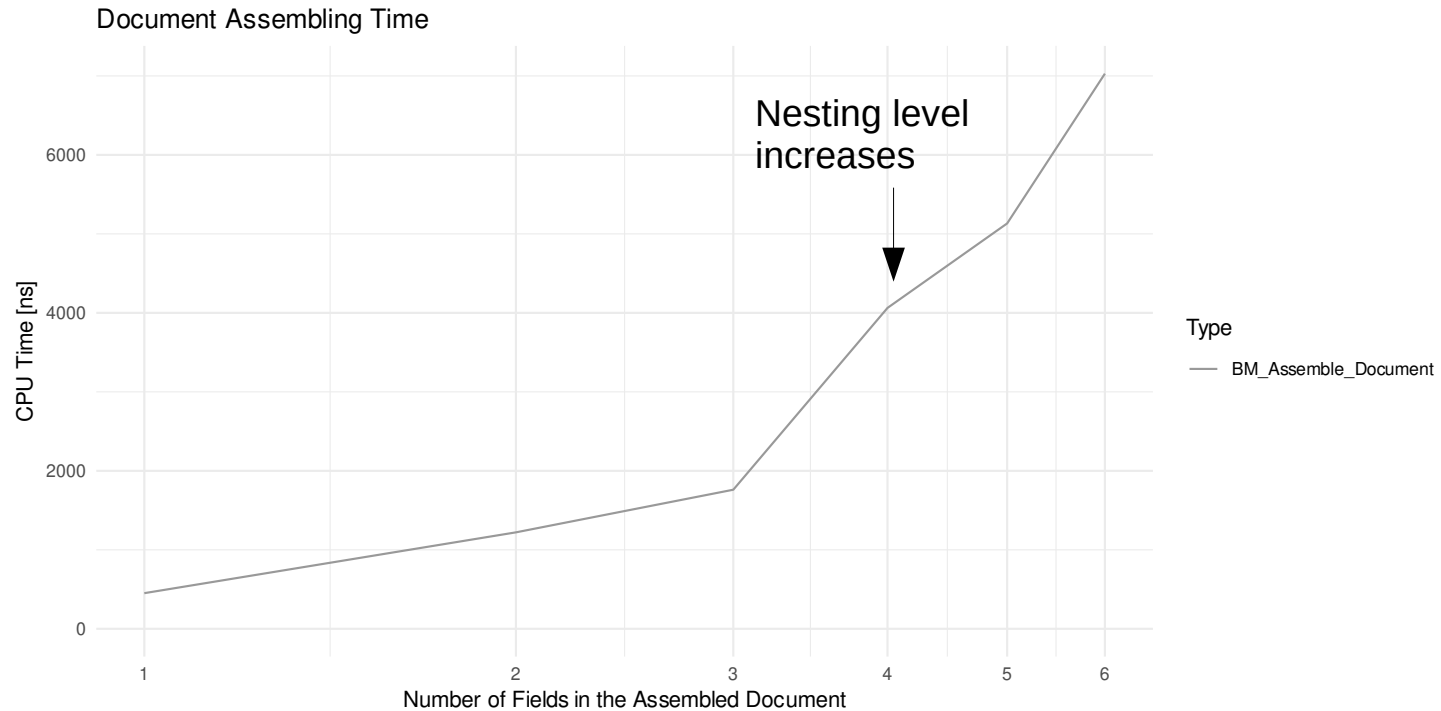
CoAn of Country and Code -> Name

Benchmarks

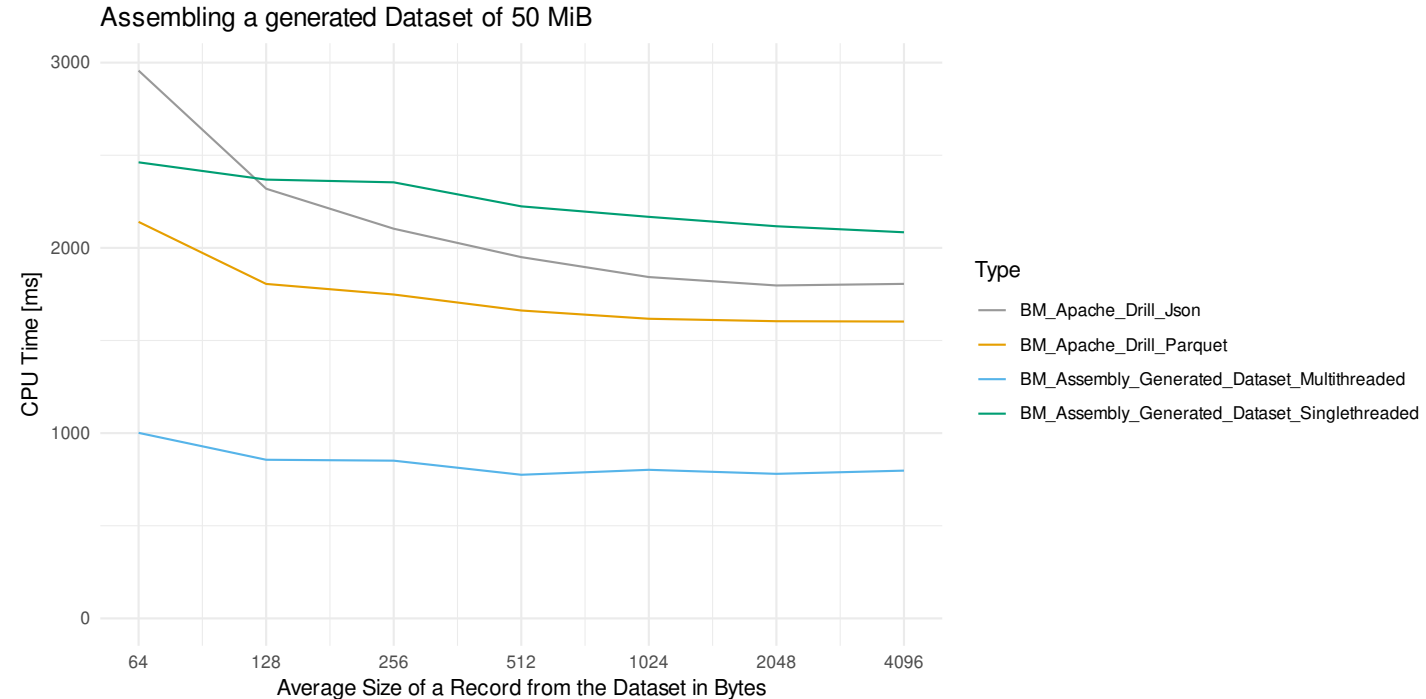
Benchmarks




Benchmarks



Performance against Apache Drill



Profiling



Time spent with Protobuf

Summary & Outlook

- FSM construction is fast and only needed once (=> neglectable overhead)
- Record assembly takes a lot of time (copies & heap allocs with Proto)
- Compilation instead of interpretation
- Not using Protobuf